Adaptive Hybrid Retrieval: Balancing Parametric and Non-Parametric Memory under Uncertainty

Authors: Dr. Tariq Al-Mansoor, Dr. Rachel Green, Dr. Liam Chen
Institution: Zurich AI Research & ETH Zurich | **Year:** 2025

ABSTRACT

The dichotomy between purely parametric generation and unconstrained RAG forces practitioners to trade off latency, cost, and hallucination rates. We introduce Adaptive Hybrid Retrieval (AHR), a dynamic entropy-gated framework that routes queries between internal parametric reasoning, dense vector search, and structured graph traversal based on model prediction uncertainty. Dynamic entropy gating routes queries to internal parametric memory when confidence is > 0.92, invoking hybrid dense-graph retrieval only on high-uncertainty prompts. This hybrid gating cuts overall query latency by 58% while maintaining 91.2% factual correctness, preventing context poisoning observed in unconstrained RAG pipelines.

1. Introduction

Neither pure RAG (Vance et al., Mehta et al.) nor pure parametric reasoning (Reynolds & Zhang) offers an optimal Pareto frontier across latency and factuality. A system must intelligently assess its own epistemic uncertainty before retrieving external context.

2. Architecture & Implementation

Our system monitors token entropy and attention dispersion across prompt representations. When epistemic uncertainty exceeds tau = 0.08, the router triggers a two-tier retrieval: fast vector retrieval for entity lookup and selective subgraph expansion for multi-hop queries.

3. Experimental Evaluation & Claims

- **Finding 1: Dynamic entropy gating routes queries to internal parametric memory when confidence is > 0.92, invoking hybrid dense-graph retrieval only on high-uncertainty prompts.**
- **Finding 2: Hybrid gating cuts overall query latency by 58% while maintaining 91.2% factual correctness across diverse benchmark suites.**
- **Finding 3: Selective retrieval prevents context poisoning observed in unconstrained RAG pipelines by suppressing retrieval when internal consensus is robust.**

4. Limitations & Research Gaps

Uncertainty calibration requires fine-tuning on domain-specific validation splits. How to adapt uncertainty thresholds in out-of-distribution clinical scenarios remains an open research challenge. Cross-institution privacy constraints were not evaluated.

5. Conclusion

Adaptive routing reconciles the conflict between retrieval latency and parametric hallucination, offering a unified path for resilient clinical LLM deployments.